

BDA Suggested Practical List:

1. To create Employee data with salary in a single table using R.

```
# 1 To create Employee data with salary in a single table using R.
# Create a data frame with employee data
employee_data <- data.frame(
  EmployeeID = c(1, 2, 3, 4, 5),
  FirstName = c("John", "Jane", "Bob", "Alice", "Charlie"),
  LastName = c("Doe", "Smith", "Johnson", "Brown", "Wilson"),
  Department = c("HR", "IT", "Sales", "Finance", "Marketing"),
  Salary = c(55000, 60000, 48000, 70000, 62000)
)
# Print the employee data
print(employee_data)
write.csv(employee_data, "employee_data.csv", row.names = FALSE)
```

2. Define a function in a script, source it to the RStudio and print the multiples of a number.

```
# Define a function to print multiples of a number
print_multiples <- function(number, n) {
  multiples <- number * 1:n
  cat("Multiples of", number, "up to", n, ":\n")
  cat(multiples, sep = ", ")
}
# Save this script as "print_multiples.R" in your working directory.
# Source the script in RStudio to make the function available in your current R session.
# Replace "D:/R_Function.txt" with the actual file path using forward slashes.
source("D:/R_Function.txt")
# Use the function to print multiples.
# For example, print multiples of 4 up to 8.
print_multiples(4, 8)
```

3. Create two data frames roll no, name, class and other data frame as roll no, subject, marks. Merge this two data frame and print the o/p.

3 Create two data frames roll no, name, class and other data frame as roll no, subject, marks.

```
# Merge this two data frame and print the o/p.
# Create the first data frame with student information
student_df <- data.frame(
  RollNo = c(1, 2, 3, 4, 5),
  Name = c("Alice", "Bob", "Charlie", "David", "Eve"),
  Class = c("A", "B", "A", "C", "B")
)
# Create the second data frame with subject and marks information
marks_df <- data.frame(
```

```

RollNo = c(2, 3, 1, 4, 5),
Subject = c("Math", "Science", "Math", "History", "English"),
Marks = c(95, 88, 92, 75, 80)
)
# Merge the two data frames based on the common column "RollNo"
merged_df <- merge(student_df, marks_df, by = "RollNo")
# Print the merged data frame
print(student_df)
print(marks_df)
print(merged_df)

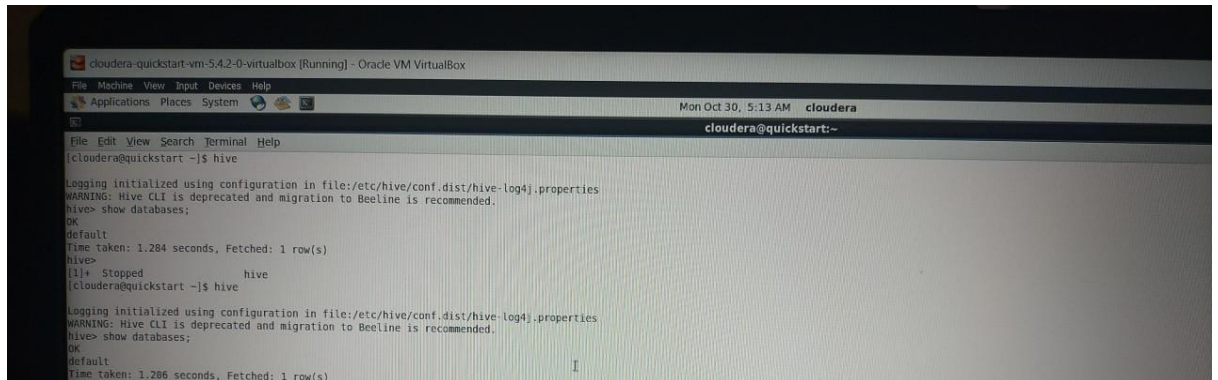
```

4. Create a table in Hive, display the employees record after loading the data into table.

```

hive
show databases;

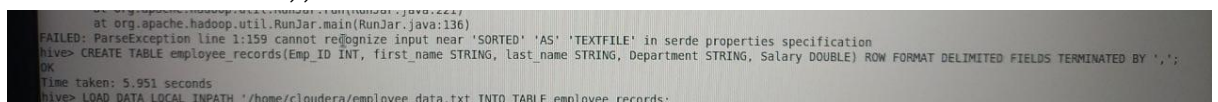
```



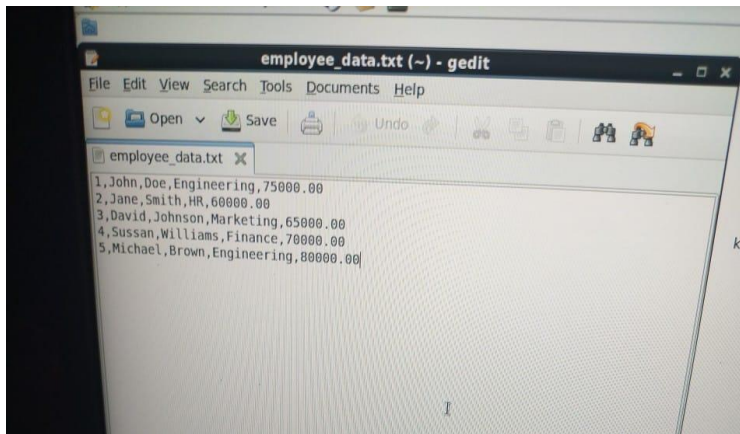
```

CREATE TABLE employee_records (
  employee_id INT,
  first_name STRING,
  last_name STRING,
  department STRING,
  salary DOUBLE
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';


```




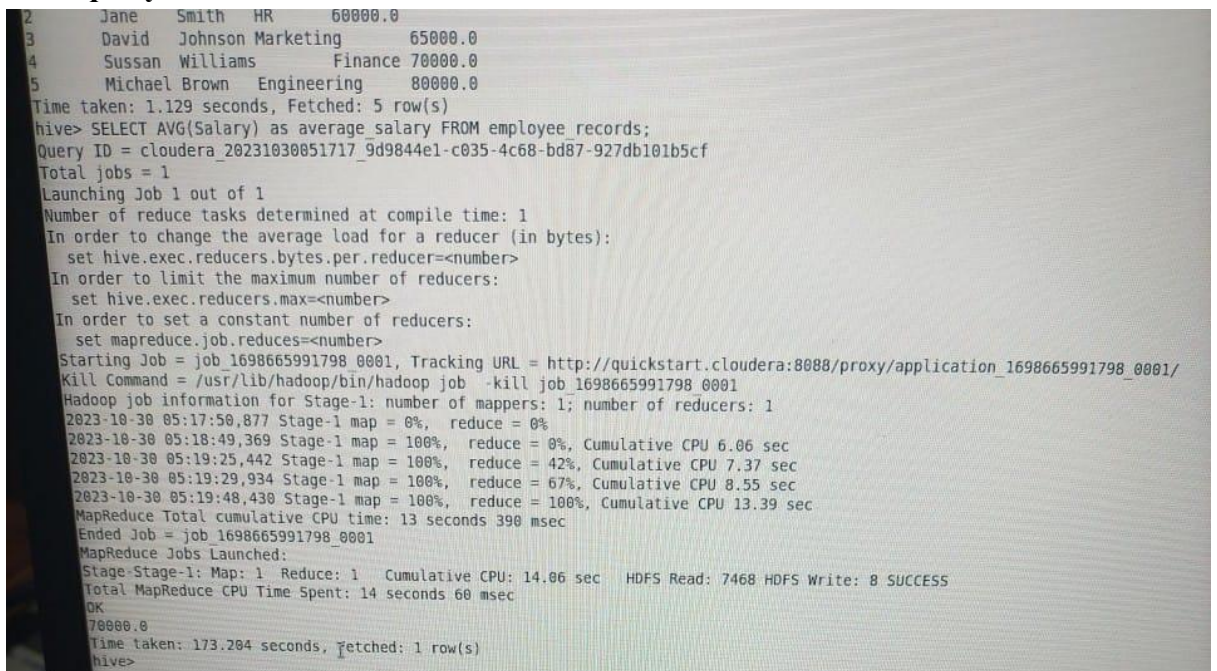
employee_data.txt



LOAD DATA LOCAL INPATH 'employee_data.txt' INTO TABLE employee_records;
SELECT *FROM employee_records;



extra query



5. Create a table in HIVE and add a partition as country='AUS'.

METHOD 1

```
CREATE TABLE employee_records (  
    employee_id INT,  
    first_name STRING,  
    last_name STRING,  
    department STRING,  
    salary DOUBLE  
)PARTITIONED BY (country STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ',';
```

```
desc employee_records;
```

```
LOAD DATA LOCAL INPATH 'employee_data_AUS.txt' OVERWRITE INTO TABLE Employee_Info  
PARTITION(country = 'AUS');
```

```
LOAD DATA LOCAL INPATH 'employee_data_US.txt' OVERWRITE INTO TABLE Employee_Info  
PARTITION(country = 'US');
```

```
SHOW PARTITIONS employee_records;
```

METHOD 2

```
CREATE TABLE employee_records (  
    employee_id INT,  
    first_name STRING,  
    last_name STRING,  
    department STRING,  
    salary DOUBLE  
)PARTITIONED BY (country STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ',';
```

```
desc Employee_Info;
```

```
LOAD DATA LOCAL INPATH 'employee_data.txt' INTO TABLE Employee_Info  
PARTITION(country = 'AUS');
```

```
SELECT *FROM Employee_Info;
```

```
ALTER TABLE employee_info ADD PARTITION (country='US');
```



```
Machine View Input Devices Help
Applications Places System
Mon Oct 30, 6:54 AM cloudera
cloudera@quickstart:~

Edit View Search Terminal Help
at java.lang.reflect.Method.invoke(Method.java:606)
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
FAILED: ParseException line 1:173 mismatched input 'AS' expecting BY near 'TERMINATED' in table row format's field separator
e> CREATE TABLE Employee_Info(Id INT,First_Name STRING,Last_Name STRING,Department STRING, Salary DOUBLE)PARTITIONED BY (country STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';

Time taken: 0.176 seconds
e> desc Employee_Info;

+-----+
| Id      |
+-----+
| First_Name |
+-----+
| Last_Name  |
+-----+
| Department |
+-----+
| Salary     |
+-----+
| Country    |
+-----+

Partition Information
+-----+
| col_name | data_type | comment |
+-----+
| Id        | int       |         |
+-----+
| First_Name | string    |         |
+-----+
| Last_Name  | string    |         |
+-----+
| Department | string    |         |
+-----+
| Salary     | double    |         |
+-----+
| Country    | string    |         |
+-----+

Partition Information
+-----+
| col_name | data_type | comment |
+-----+
| Id        | int       |         |
+-----+
| First_Name | string    |         |
+-----+
| Last_Name  | string    |         |
+-----+
| Department | string    |         |
+-----+
| Salary     | double    |         |
+-----+
| Country    | string    |         |
+-----+

Time taken: 0.225 seconds, Fetched: 11 row(s)
hive> LOAD DATA LOCAL INPATH 'employee_data.txt' INTO TABLE Employee_Info;
FAILED: SemanticException [Error 10002]: Need to specify partition columns because the destination table is partitioned
hive> LOAD DATA LOCAL INPATH 'employee_data.txt' INTO TABLE Employee_Info PARTITION (country = 'AUS');
Loading data to table default.employee_info partition (country=AUS)
Partition default.employee_info(country=AUS) stats: [numFiles=1, numRows=0, totalSize=164, rawDataSize=0]
OK
Time taken: 2.646 seconds
hive> SELECT *FROM Employee_Info;

+-----+
| Id      |
+-----+
| First_Name |
+-----+
| Last_Name  |
+-----+
| Department |
+-----+
| Salary     |
+-----+
| Country    |
+-----+
| Id      | First_Name | Last_Name | Department | Salary | Country |
+-----+
| 1       | John       | Doe       | Engineering | 75000.0 | AUS     |
+-----+
| 2       | Jane       | Smith     | HR          | 60000.0 | AUS     |
+-----+
| 3       | David      | Johnson   | Marketing   | 65000.0 | AUS     |
+-----+
| 4       | Sussan     | Williams  | Finance     | 70000.0 | AUS     |
+-----+
| 5       | Michael    | Brown     | Engineering | 80000.0 | AUS     |
+-----+
Time taken: 0.47 seconds, Fetched: 5 row(s)
hive> ALTER TABLE Employee_Info ADD PARTITION (country = 'USA');
Time taken: 0.314 seconds
hive> SELECT *FROM Employee_Info;

+-----+
| Id      |
+-----+
| First_Name |
+-----+
| Last_Name  |
+-----+
| Department |
+-----+
| Salary     |
+-----+
| Country    |
+-----+
| Id      | First_Name | Last_Name | Department | Salary | Country |
+-----+
| 1       | John       | Doe       | Engineering | 75000.0 | AUS     |
+-----+
| 2       | Jane       | Smith     | HR          | 60000.0 | AUS     |
+-----+
| 3       | David      | Johnson   | Marketing   | 65000.0 | AUS     |
+-----+
| 4       | Sussan     | Williams  | Finance     | 70000.0 | AUS     |
+-----+
| 5       | Michael    | Brown     | Engineering | 80000.0 | AUS     |
+-----+
Time taken: 0.283 seconds, Fetched: 5 row(s)
hive> SHOW PARTITIONS;
NoViableAltException(-10[184:1: tableName : (db= identifier DOT tab= identifier -> ^( TOK_TABNAME $db $tab) |tab= identifier -> ^( TOK_TABNAME $tab) :);])
```

SHOW PARTITIONS employee_info;

```
Machine View Input Devices Help
Applications Places System
Mon Oct 30, 6:54 AM cloudera
cloudera@quickstart:~

Edit View Search Terminal Help
Time taken: 0.47 seconds, Fetched: 5 row(s)
hive> ALTER TABLE Employee_Info ADD PARTITION (country = 'USA');
OK
Time taken: 0.314 seconds
hive> SELECT *FROM Employee_Info;

+-----+
| Id      |
+-----+
| First_Name |
+-----+
| Last_Name  |
+-----+
| Department |
+-----+
| Salary     |
+-----+
| Country    |
+-----+
| Id      | First_Name | Last_Name | Department | Salary | Country |
+-----+
| 1       | John       | Doe       | Engineering | 75000.0 | AUS     |
+-----+
| 2       | Jane       | Smith     | HR          | 60000.0 | AUS     |
+-----+
| 3       | David      | Johnson   | Marketing   | 65000.0 | AUS     |
+-----+
| 4       | Sussan     | Williams  | Finance     | 70000.0 | AUS     |
+-----+
| 5       | Michael    | Brown     | Engineering | 80000.0 | AUS     |
+-----+
Time taken: 0.283 seconds, Fetched: 5 row(s)
hive> SHOW PARTITIONS;
NoViableAltException(-10[184:1: tableName : (db= identifier DOT tab= identifier -> ^( TOK_TABNAME $db $tab) |tab= identifier -> ^( TOK_TABNAME $tab) :);])
at org.apache.hadoop.hive.ql.parse.HiveParser.FromClauseParser.tableName(HiveParser.FromClauseParser.java:4674)
at org.apache.hadoop.hive.ql.parse.HiveParser.tableName(HiveParser.java:44373)
at org.apache.hadoop.hive.ql.parse.HiveParser.showStatement(HiveParser.java:28856)
at org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:2430)
at org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:1579)
at org.apache.hadoop.hive.ql.parse.HiveParser.parseStatement(HiveParser.java:1057)
at org.apache.hadoop.hive.ql.parse.ParserDriver.parse(ParseDriver.java:199)
at org.apache.hadoop.hive.ql.parse.ParserDriver.parse(ParseDriver.java:166)
at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:393)
at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:307)
at org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1110)
at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1047)
at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1037)
at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:207)
at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:159)
at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:376)
at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:756)
at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:675)
at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:615)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:606)
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
FAILED: ParseException line 1:15 cannot recognize input near '<EOF>' '<EOF>' in table name
hive> SHOW PARTITIONS Employee_Info;
OK
country=AUS
country=USA
Time taken: 0.174 seconds, Fetched: 2 row(s)
hive>
```

6. Hadoop Basic commands –

Hdfs dfs -ls /

Hdfs dfs -mkdir /user/cloudera/swarali

Hdfs dfs -touchz /home/cloudera/sample.txt

Hdfs dfs -mkdir /user/cloudera/swarali/

```

Hdfs dfs -copyFromLocal /home/cloudera/a1.txt /user/cloudera/swarali
Hdfs dfs -copyToLocal /home/cloudera/swarali/sample.txt /home/cloudera
Hdfs dfs -cat /user/cloudera/swarali/a1.txt
Hdfs dfs -rm /user/cloudera/swarali/a1.txt
Hdfs dfs -rmdir /user/cloudera/swarali
Hdfs dfs -cp /user/cloudera/a1.txt /user/cloudera/swarali/
Hdfs dfs -mv /user/cloudera/a1.txt /user/cloudera/swarali/
Hdfs dfs -chmod 750 /user/cloudera/swarali/a1.txt
Hdfs dfs -chown swarali /user/cloudera/swarali/a1.txt
Hdfs dfs -chgrp swarali /user/cloudera/swarali/a1.txt
Hdfs dfs -getmerge /user/cloudera/swarali /home/cloudera/merged.txt
Hdfs dfs -put /home/cloudera/merged.txt /user/cloudera/swarali
Hdfs dfs -get /user/cloudera/sample.txt /home/cloudera
Hdfs dfs -du /user/cloudera/swarali/sample.txt
Hdfs dfsadmin -safemode enter
Hdfs dfsadmin -safemode leave
Hdfs dfsadmin -refreshNodes

```

7. Create and insert a post in NOSQL database for a social media website using MongoDB.

use your_database_name

```
db.createCollection("your_collection_name")
```

```

db.your_collection_name.insertOne({
  title: "My First Post",
  content: "This is the content of my first post.",
  author: "John Doe",
  timestamp: new Date()
})

```

```
db.your_collection_name.find()
```

8. Create and update a post in NOSQL database for a social media website using MongoDB

use your_database_name

```
db.createCollection("your_collection_name")
```

```
db.your_collection_name.insertOne({
    title: "My First Post",
    content: "This is the content of my first post.",
    author: "John Doe",
    timestamp: new Date()
})
```

```
db.your_collection_name.find().pretty()
```

```
db.your_collection_name.updateOne(
    { title: "My First Post" },
    { $set: { content: "Updated content" } }
)
```

```
db.your_collection_name.find().pretty()
```

9. Implement FM algorithm

```
import time
#BE3_22_SwaraliJanaskar_Expt-6 import time
def flajolet_martin(stream):
    maxnum = 0
    for i in range(0, len(stream)):
        val = bin((1 * stream[i] + 6) % 32)[2:]
        sum = 0
        for j in range(len(val) - 1, 0, -1):
            if val[j] == '0':
                sum += 1
            else:
                break
        if sum > maxnum:
            maxnum = sum
    return 2 ** maxnum
if __name__ == "__main__":
    stream = [1, 2, 3, 4, 5, 6, 4, 2, 5, 9, 1, 6, 3, 7, 1, 2, 2, 4, 2,
1]
    print('Using Flajolet Martin Algorithm:')
    start_time = time.time()
```

```
distinct_elements = flajolet_martin(stream)
print('Distinct elements:', distinct_elements)
print('Execution time:', time.time() - start_time, 'seconds')
```

The code you've provided is implementing the Flajolet-Martin algorithm to estimate the number of distinct elements in a stream of data. The algorithm is a probabilistic algorithm used for cardinality estimation and is particularly useful when you have a large dataset and you want to estimate the number of unique elements without storing all of them in memory.

Here's how the Flajolet-Martin algorithm works:

- It processes the elements in the stream one by one.
- For each element in the stream, it computes a binary representation of the element. In this code, the binary representation is computed as `bin((1 * stream[i] + 6) % 32)[2:]`.
- It counts the number of trailing zeros in the binary representation (i.e., the number of consecutive '0' bits from the right).
- It keeps track of the maximum count of trailing zeros encountered in the binary representations of the elements.
- The algorithm returns an estimate of the number of distinct elements as 2 raised to the power of the maximum count of trailing zeros.

Certainly, let's break down the code step by step:

```
```python
def flajolet_martin(stream):
 maxnum = 0 # Initialize the maximum number of trailing zeros to 0
 for i in range(0, len(stream)):
 val = bin((1 * stream[i] + 6) % 32)[2:]
    ```
```

1. `maxnum` is initialized to 0, which will keep track of the maximum count of trailing zeros in the binary representations of elements in the stream.

2. The code iterates through the `stream`, which is a list of numbers.

3. For each element in the stream, it computes a binary representation. It does this by taking the element, multiplying it by 1, adding 6, and taking the result modulo 32. This ensures that the number is in the range [0, 31]. Then, `bin()` is used to obtain the binary representation of the number, and `[2:]` is used to remove the '0b' prefix from the binary string.

```
```python
 sum = 0
 for j in range(len(val) - 1, 0, -1):
 if val[j] == '0':
 sum += 1
 else:
 break
    ```
```



```
'''
```

4. For each binary representation, the code initializes `sum` to 0 and starts iterating through the binary string from right to left (from the least significant bit to the most significant bit).

5. It counts the number of trailing zeros (consecutive '0' bits) in the binary representation. If it encounters a '1', the loop breaks because it's no longer counting trailing zeros.

```
'''python
    if sum > maxnum:
        maxnum = sum
'''
```

6. It compares the `sum` (the count of trailing zeros) with the current maximum count (`maxnum`). If the new `sum` is greater than the current `maxnum`, it updates `maxnum` with the new value.

```
'''python
    return 2 ** maxnum
'''
```

7. After processing all elements in the stream, the function returns an estimate of the number of distinct elements. It does this by taking 2 raised to the power of `maxnum`. The idea is that the maximum count of trailing zeros in the binary representations can be used to estimate the number of distinct elements probabilistically.

8. In the main part of the code, a sample `stream` is provided, and the Flajolet-Martin algorithm is applied to estimate the number of distinct elements in the stream. The execution time is also measured and printed.

The Flajolet-Martin algorithm is a probabilistic algorithm and provides an approximate estimate of the number of distinct elements in a large dataset, which can be useful in various applications like estimating unique website visitors or counting unique items in a dataset without having to store all the unique items explicitly. The accuracy of the estimate depends on the characteristics of the data and the chosen hash functions.

10. Implement Blooms Filter and verify the contents are present or not.

```
import mmh3
from bitarray import bitarray

class BloomFilter:
    def __init__(self, size, hash_functions):
        self.size = size
        self.bit_array = bitarray(size)
        self.bit_array.setall(0)
        self.hash_functions = hash_functions

    def add(self, element):
        for i in range(self.hash_functions):
            index = mmh3.hash(element, i) % self.size
            self.bit_array[index] = 1
```

```

def contains(self, element):
    for i in range(self.hash_functions):
        index = mmh3.hash(element, i) % self.size
        if not self.bit_array[index]:
            return False
    return True

if __name__ == "__main__":
    size = 1000
    hash_functions = 5
    bloom_filter = BloomFilter(size, hash_functions)

    # Add elements to the Bloom Filter
    elements_to_add = ["apple", "banana", "cherry", "date",
"elderberry"]
    for element in elements_to_add:
        bloom_filter.add(element)

    # Verify if elements are present
    elements_to_verify = ["apple", "banana", "grape", "fig"]
    for element in elements_to_verify:
        if bloom_filter.contains(element):
            print(f'{element}' may be in the set.")
        else:
            print(f'{element}' is definitely not in the set.")

```

Certainly, let's go through each line of the provided code and explain its purpose:

```

```python
import mmh3
from bitarray import bitarray
```

```

- These lines import the necessary libraries. `mmh3` is used for hash functions, and `bitarray` is used for working with a bit array.

```

```python
class BloomFilter:
 def __init__(self, size, hash_functions):
 self.size = size
 self.bit_array = bitarray(size)
 self.bit_array.setall(0)
 self.hash_functions = hash_functions
```

```

- This code defines a class named `BloomFilter`. The class has an `__init__` method that initializes a new instance of the BloomFilter with the specified size and number of hash functions.
- `self.size` stores the size of the bit array, which is passed as a parameter to the constructor.
- `self.bit_array` is a bit array of the specified size created using the `bitarray` library. It's initialized with all bits set to 0 using `setall(0)`.
- `self.hash_functions` stores the number of hash functions to be used.

```

```python
def add(self, element):
 for i in range(self.hash_functions):
 index = mmh3.hash(element, i) % self.size
 self.bit_array[index] = 1
...

```

- The `add` method is used to add an element to the Bloom Filter.
- It iterates through each of the specified hash functions (from 0 to `hash\_functions - 1`).
- For each hash function, it calculates an index within the bit array by taking the result of `mmh3.hash(element, i)` modulo the size of the bit array (`self.size`).
- It then sets the corresponding bit at that index to 1, indicating that the element is present in the set.

```

```python
def contains(self, element):
    for i in range(self.hash_functions):
        index = mmh3.hash(element, i) % self.size
        if not self.bit_array[index]:
            return False
    return True
...

```

- The `contains` method is used to check if an element may be in the set (with a possibility of false positives).
- It also iterates through each of the specified hash functions.
- For each hash function, it calculates the index within the bit array as before.
- If any of the bits at the calculated indexes are not set (equal to 0), it returns `False` because the element is definitely not in the set.
- If all bits at the calculated indexes are set (equal to 1), it returns `True`, indicating that the element may be in the set.

```

```python
if __name__ == "__main__":
 size = 1000
 hash_functions = 5
 bloom_filter = BloomFilter(size, hash_functions)

 # Add elements to the Bloom Filter
 elements_to_add = ["apple", "banana", "cherry", "date", "elderberry"]
 for element in elements_to_add:
 bloom_filter.add(element)

 # Verify if elements are present
 elements_to_verify = ["apple", "banana", "grape", "fig"]
 for element in elements_to_verify:

```

```

 if bloom_filter.contains(element):
 print(f"{element}' may be in the set.")
 else:
 print(f"{element}' is definitely not in the set.")
 ...

```

- This section is the main block of the program.
- It initializes a Bloom Filter instance `bloom\_filter` with a bit array of size 1000 and 5 hash functions.
- It adds several elements to the Bloom Filter using the `add` method.
- It then verifies if certain elements are present in the Bloom Filter using the `contains` method. For each element, it prints whether the element may be in the set or is definitely not in the set based on the results of the Bloom Filter.

Overall, this code demonstrates the use of a Bloom Filter to test for potential set membership with possible false positives.

## 11. Implement Word count using map reduce.

**Step 1 :** New->Java project -> Project\_name

Add external libraries

Libraies->Add External jar files -> file system -> user ->select all jar files->

finish

**Step 2 :** Right click on src-> new-> class->class\_name(WordCount)->next

**Step 3 :** copy code from the above link and save the file

<https://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

Code:

```

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

 public static class TokenizerMapper
 extends Mapper<Object, Text, Text, IntWritable>{

 private final static IntWritable one = new IntWritable(1);
 private Text word = new Text();

 public void map(Object key, Text value, Context context
) throws IOException, InterruptedException {

```

```

StringTokenizer itr = new StringTokenizer(value.toString());
while (itr.hasMoreTokens()) {
 word.set(itr.nextToken());
 context.write(word, one);
}
}
}

```

```

public static class IntSumReducer
 extends Reducer<Text,IntWritable,Text,IntWritable> {
 private IntWritable result = new IntWritable();

 public void reduce(Text key, Iterable<IntWritable> values,
 Context context
) throws IOException, InterruptedException {
 int sum = 0;
 for (IntWritable val : values) {
 sum += val.get();
 }
 result.set(sum);
 context.write(key, result);
 }
}

```

```

public static void main(String[] args) throws Exception {
 Configuration conf = new Configuration();
 Job job = Job.getInstance(conf, "word count");
 job.setJarByClass(WordCount.class);
 job.setMapperClass(TokenizerMapper.class);
 job.setCombinerClass(IntSumReducer.class);
 job.setReducerClass(IntSumReducer.class);
 job.setOutputKeyClass(Text.class);
 job.setOutputValueClass(IntWritable.class);
 FileInputFormat.addInputPath(job, new Path(args[0]));
 FileOutputFormat.setOutputPath(job, new Path(args[1]));
 System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

**Step 4 :** click on export->Java->jar file->next->browse->save jar file as WordCount.jar->location choose desktop->next->finish

**Step 5 :** Create Sample.txt and save it on desktop

**Step 6 :** Create directory


```
hadoop dfs -mkdir /WordIn
```

```
cd Desktop/
```

```
hadoop dfs -copyFromLocal Sample.txt /WordIn
```

```
hadoop jar WordCount.jar WordCount /WordIn /WordOut
```





**12. Write a R program to create a Data frames which contain details of 5 employees and display the details.**

```
12 Write a R program to create a Data frames
which contain details of 5 employees and display the details.
Create a data frame with employee details
employee_df <- data.frame(
 EmployeeID = c(1, 2, 3, 4, 5),
 FirstName = c("John", "Jane", "Bob", "Alice", "Charlie"),
 LastName = c("Doe", "Smith", "Johnson", "Brown", "Wilson"),
 Department = c("HR", "IT", "Sales", "Finance", "Marketing"),
 Salary = c(55000, 60000, 48000, 70000, 62000)
)
Display the employee details
print(employee_df)
```

**13. Write a R program to create a list of heterogeneous data, which include character, numeric and logical vectors. Print the lists.**

```
13 Write a R program to create a list of heterogeneous data,
which include character, numeric and logical vectors. Print the lists.
Create a list of heterogeneous data
my_list <- list(
 character_vector = c("Alice", "Bob", "Charlie"),
 numeric_vector = c(25, 30, 35),
 logical_vector = c(TRUE, FALSE, TRUE)
)
```

```
Print the list
print(my_list)
```

**14. Write a R program to create bell curve of a random normal distribution.**

```
14 Write a R program to create bell curve of a random normal distribution.
Set the seed for reproducibility (optional)
set.seed(123)
Generate random data from a normal distribution
data <- rnorm(1000, mean = 0, sd = 1)
Create a histogram to visualize the distribution
hist(data, main = "Random Normal Distribution", xlab = "Value", ylab = "Frequency", col =
"lightblue", border = "black", breaks = 20)
when data is given
Create a vector with your existing data
your_data <- c(0.5, 1.2, 1.8, 2.4, 2.9, 3.2, 3.5, 4.0, 4.5, 5.0, 5.5, 5.8, 6.2, 6.5, 6.8)
Create a histogram to visualize the data
hist(your_data, main = "Your Data Distribution", xlab = "Value", ylab = "Frequency", col =
"lightblue", border = "black", breaks = 10)
```

**15. Write a R program to create a simple bar plot of five subjects marks.**

```
15 Write a R program to create a simple bar plot of five subjects marks.
Create a vector of subject names
subjects <- c("Math", "Science", "English", "History", "Art")
Create a vector of corresponding marks
marks <- c(90, 78, 85, 92, 70)
Create a bar plot
barplot(marks, names.arg = subjects, col = "skyblue", main = "Subject Marks", xlab =
"Subjects", ylab = "Marks")
```

**16. Write a R program to create a  $5 \times 4$  matrix ,  $3 \times 3$  matrix with labels and fill the matrix by rows and  $2 \times 2$  matrix with labels and fill the matrix by columns.**

```
16 Write a R program to create a 5×4 matrix , 3×3 matrix with labels and fill the matrix by
rows
```

```
and 2×2 matrix with labels and fill the matrix by columns.
```

```
Create a 5x4 matrix with labels and fill it by rows
```

```
matrix1 <- matrix(1:20, nrow = 5, ncol = 4, byrow = TRUE, dimnames = list(c("Row1",
"Row2", "Row3", "Row4", "Row5"), c("Col1", "Col2", "Col3", "Col4")))
```

```
Create a 3x3 matrix with labels and fill it by rows
```

```
matrix2 <- matrix(21:29, nrow = 3, ncol = 3, byrow = TRUE, dimnames = list(c("RowA",
"RowB", "RowC"), c("ColX", "ColY", "ColZ")))
```

```

Create a 2x2 matrix with labels and fill it by columns

matrix3 <- matrix(30:33, nrow = 2, ncol = 2, byrow = FALSE, dimnames = list(c("RowI",
"RowII"), c("ColM", "ColN")))

Print the matrices

print(matrix1)

print(matrix2)

print(matrix3)

```

**17. Write a R program to create a sequence of numbers from 20 to 50 and find the mean of numbers from 20 to 60 and sum of numbers from 51 to 91.**

```

17 Write a R program to create a sequence of numbers from 20 to 50
and find the mean of numbers from 20 to 60 and sum of numbers from 51 to 91.

Create a sequence of numbers from 20 to 50
numbers_seq <- seq(20, 50)

Calculate the mean of numbers from 20 to 60
mean_20_to_60 <- mean(seq(20, 60))

Calculate the sum of numbers from 51 to 91
sum_51_to_91 <- sum(seq(51, 91))

Print the results

cat("Sequence from 20 to 50:", numbers_seq, "\n")

cat("Mean of numbers from 20 to 60:", mean_20_to_60, "\n")

cat("Sum of numbers from 51 to 91:", sum_51_to_91, "\n")

```